



Pure Storage and Open Source

Simon Dodsley

Director, Open Source Integrations
(theansibleguy.com)

@purestorage

April 30, 2024



Agenda

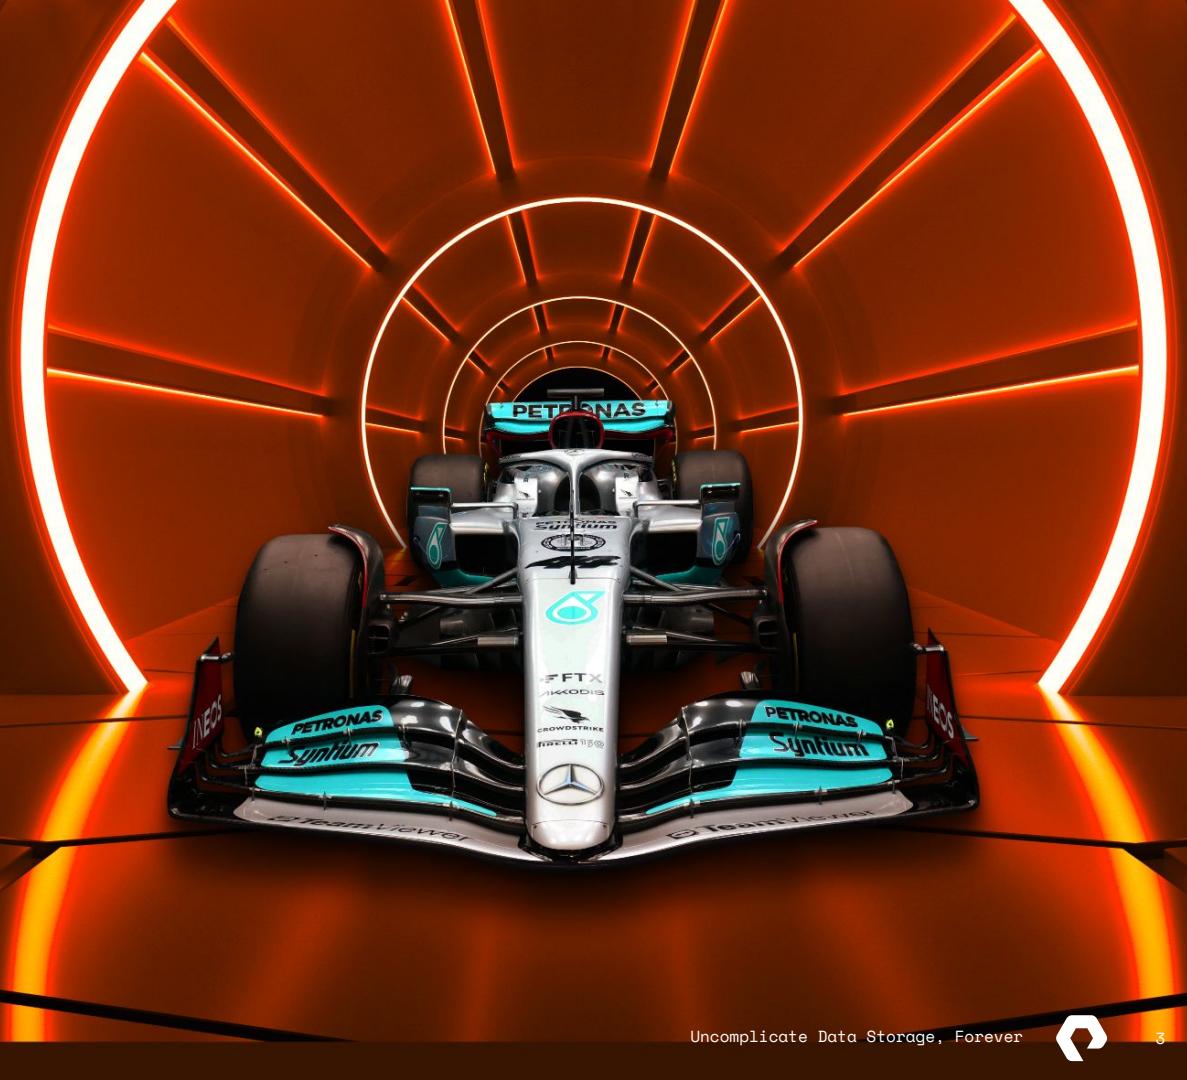
Orchestration

Automation

Observability



Orchestration





ONE OF THE MOST ACTIVE OPEN SOURCE PROJECTS WITH 40 MILLION CORES IN PRODUCTION

- OpenStack community completes the alphabet with Zed, and moves on to date-based releases
 - **2024.1 (Caracal)** just released
- OpenStack production deployments hit 40 Million Cores
- More than 450 organizations have contributed upstream to OpenStack, making it one of the most active open source projects in the world
- Over 9000 contributors across 180 countries
- OpenStack public cloud footprint exceeds 300 data centers worldwide

OpenStack

An IaaS Platform – without all those pesky licensing costs...



Cinder Block Storage
(FlashArray)



Manila File Storage
(FlashBlade)



Block QoS



Simple Management

Block Replication



iSCSI, FC, NVMe, NFS

Multi-platform



Enhanced user experience



Red Hat, Canonical, IBM,
Mirantis & Platform9

Release Cycles

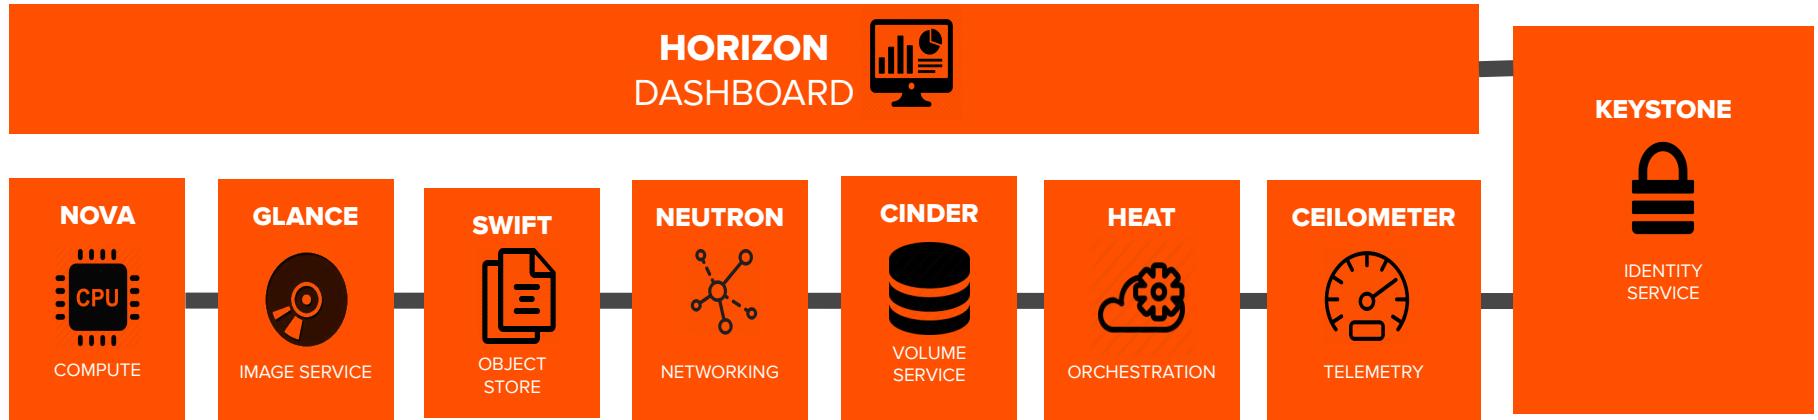
Six monthly cycle

- Releases correspond to Design Summits
- Alphabetic release naming extended to date-based
- Latest Release: 2024.1 (Caracal)
- Next Release: 2024.2 (Dalmatian)
- Upgrade paths now available, including skip-level
- Less delta between releases now
- More feature rich

Series	Status	Initial Release Date	Next Phase	EOL Date
2024.2 Dalmatian	Development	2024-10-02 <i>estimated (schedule)</i>	Maintained	<i>estimated 2024-10-02</i>
2024.1 Caracal (SLURP)	Maintained	2024-04-03	Unmaintained	<i>estimated 2025-10-03</i>
2023.2 Bobcat	Maintained	2023-10-04	End Of Life	<i>estimated 2025-04-04</i>
2023.1 Antelope (SLURP)	Maintained	2023-03-22	Unmaintained	<i>estimated 2024-09-22</i>
Zed	Maintained	2022-10-05	Unmaintained	<i>estimated 2024-05-02</i>
Yoga	Unmaintained	2022-03-30	End Of Life	TBD
Xena	Unmaintained	2021-10-06	End Of Life	TBD
Wallaby	Unmaintained	2021-04-14	End Of Life	TBD
Victoria	Unmaintained	2020-10-14	End Of Life	TBD
Ussuri	End Of Life	2020-05-13		2024-02-20
Train	End Of Life	2019-10-16		2024-02-20
Stein	End Of Life	2019-04-10		2024-01-04
Rocky	End Of Life	2018-08-30		2023-08-16
Queens	End Of Life	2018-02-28		2023-01-18
Pike	End Of Life	2017-08-30		2022-10-13
Ocata	End Of Life	2017-02-22		2021-06-15
Newton	End Of Life	2016-10-06		2017-10-25
Mitaka	End Of Life	2016-04-07		2017-04-10
Liberty	End Of Life	2015-10-15		2016-11-17
Kilo	End Of Life	2015-04-30		2016-05-02
Juno	End Of Life	2014-10-16		2015-12-07
Icehouse	End Of Life	2014-04-17		2015-07-02
Havana	End Of Life	2013-10-17		2014-09-30
Grizzly	End Of Life	2013-04-04		2014-03-29
Folsom	End Of Life	2012-09-27		2013-11-19
Essex	End Of Life	2012-04-05		2013-05-06
Diablo	End Of Life	2011-09-22		2013-05-06
Cactus	End Of Life	2011-04-15		
Bexar	End Of Life	2011-02-03		
Austin	End Of Life	2010-10-21		

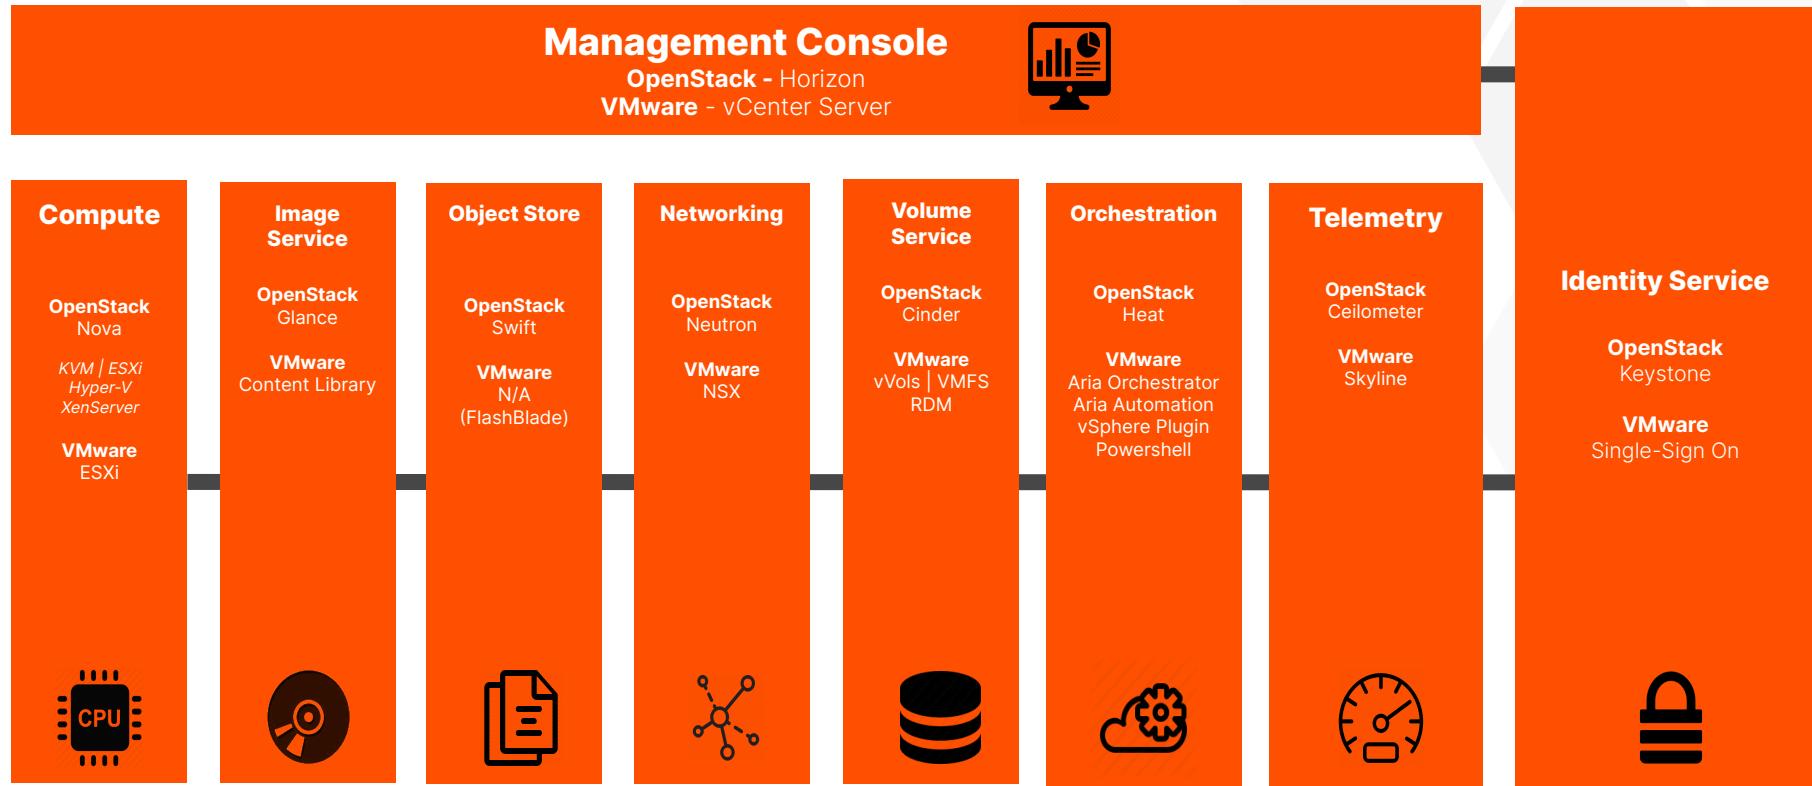


Architecture



- Modular architecture
- Designed for easy scale-out
- Based of (growing) set of core services

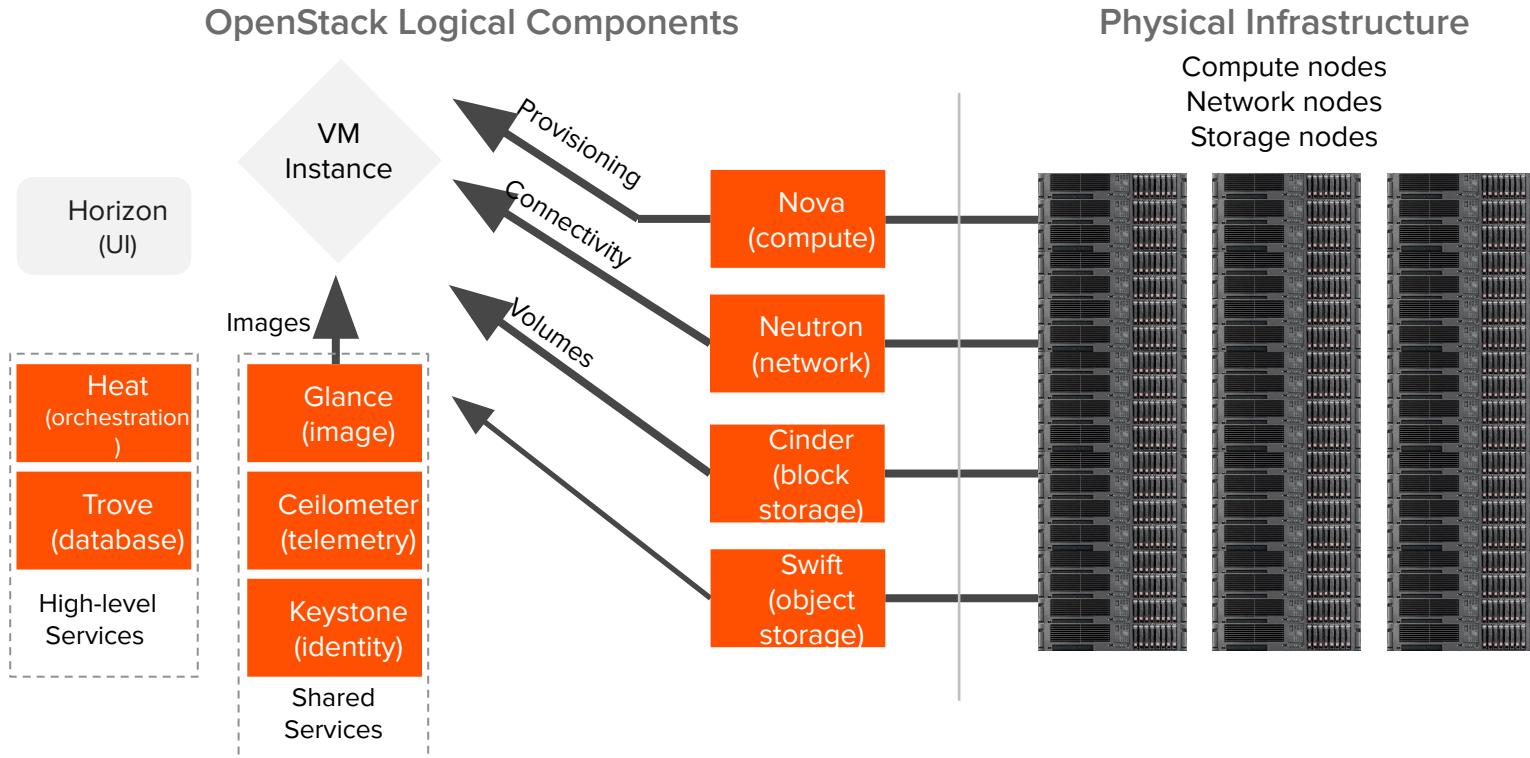
Feature Comparison - OpenStack / VMware



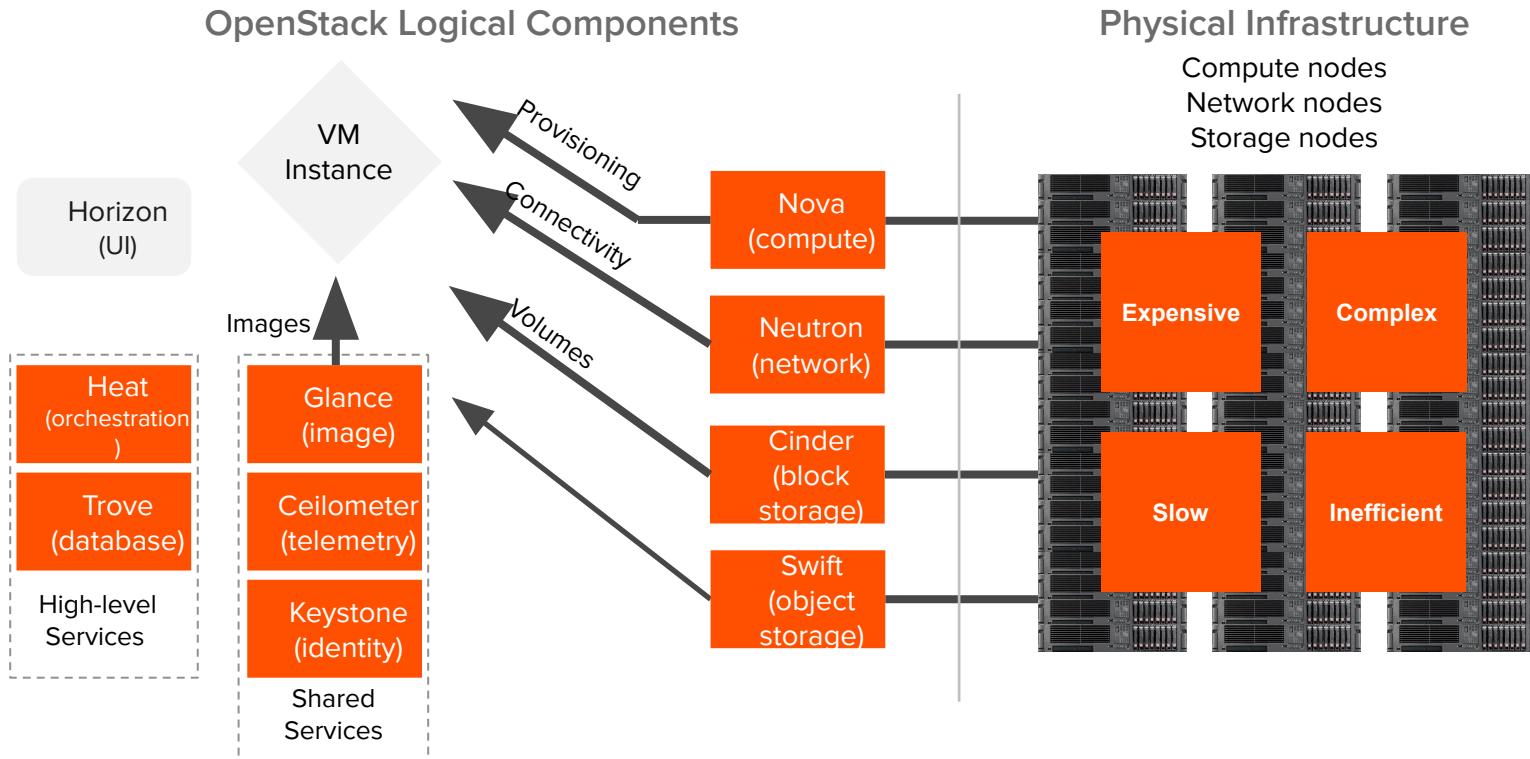
Projects

- **Important Others:**
 - Manila – Filesystem service
 - Trove – DBaaS
 - Ironic – Bare Metal
 - Magnum - Container orchestration

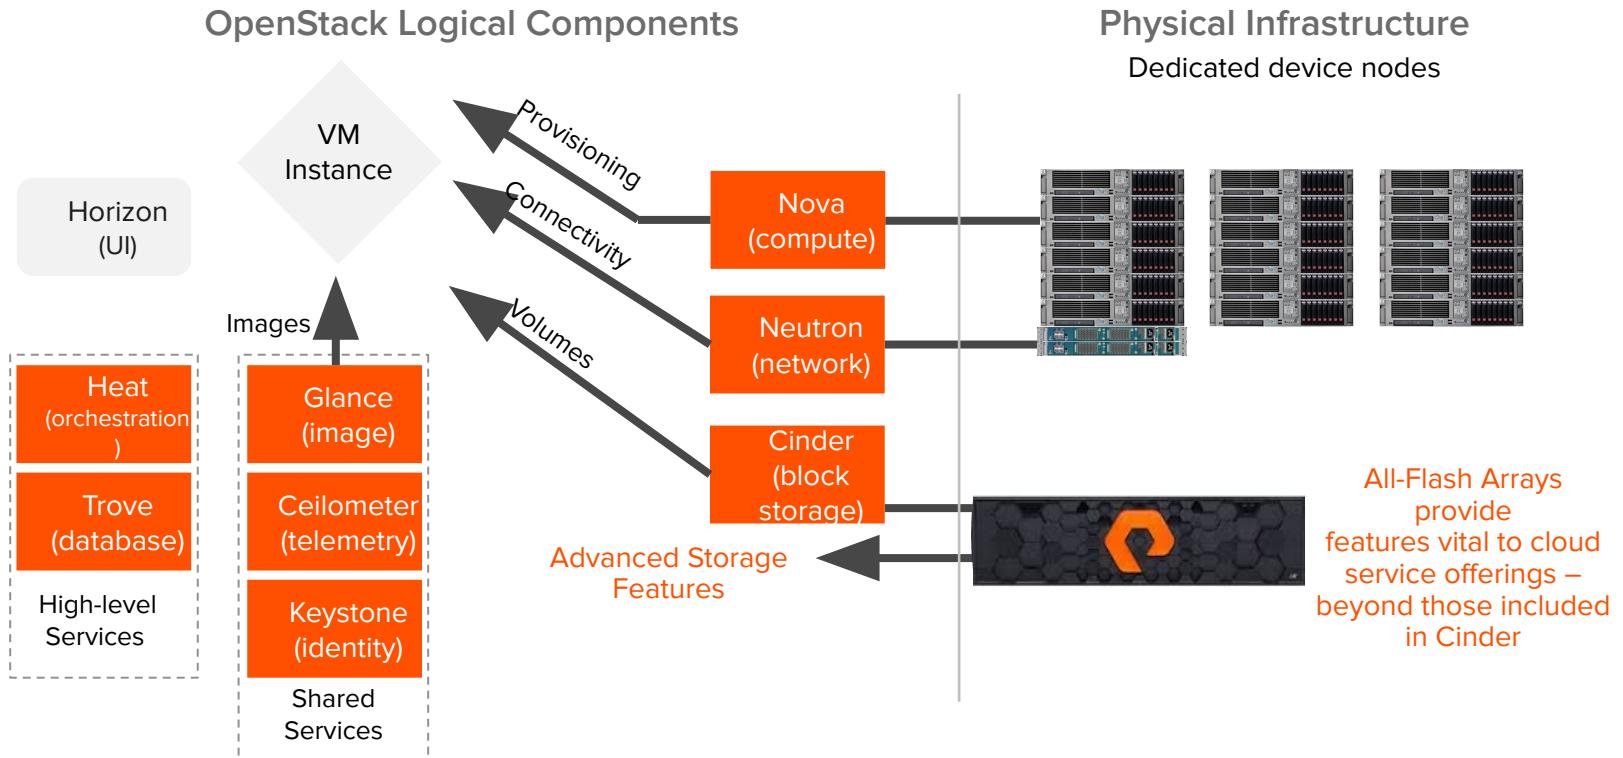
Typical OpenStack Deployment With White Box Systems



Typical OpenStack Deployment With White Box Systems



Typical OpenStack Deployment With White Box System

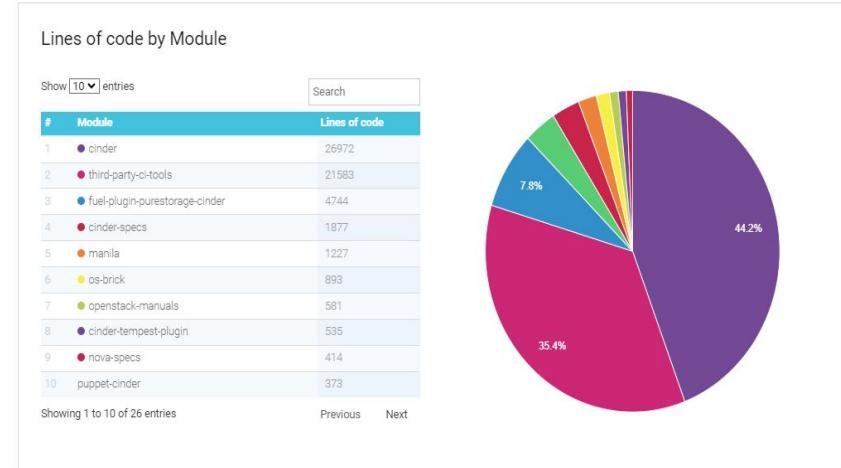




Pure and OpenStack

OpenStack Contributions

- Over 60,900 lines of code (27,000 in Cinder)
- Continuous Contributor to Cinder
- Code to 26 projects
- High-Level Architectural Contributions
- Supported Vendor Driver



Pure's Investment In OpenStack

- Full in-tree Cinder driver for FlashArray since Juno (2014 H2)
- Unique Multi-protocol Cinder support – iSCSI, FC, NVMe-RoCE & NVMe-TCP
- Full in-tree Manila driver for FlashBlade since Xena (2021 H2)
- Active contributor to multiple core projects
- Continuing development of driver – new feature add and feature support
- Certified with Red Hat, Canonical, Mirantis and IBM OpenStack deployments
- Dedicated Corporate team of developers and architects
- Sponsorship of Open Infrastructure Summits and Days



OpenStack Cinder Support Matrix

Pure Storage Related

•Required Driver Functions:

- Create Volume
- Delete Volume
- Attach Volume
- Detach Volume
- Extend Volume
- Create Snapshot
- Delete Snapshot
- Create Volume from Snapshot
- Create Volume from Volume (clone)
- Create Image from Volume
- Volume Migration (host assisted)

•Optional Driver Functions Supported:

- Extend an Attached Volume
- Volume Replication (async, sync & trisync)
- Consistency Groups
- Replicated CG Snapshots
- Thin Provisioning
- Multi-Attach Support
- Revert to Snapshot
- iSCSI & NVMe CIDR support (IPv4 & IPv6)
- iSCSI & NVMe CIDR list support (IPv4 & IPv6)
- Host Personality
- Quality-of-Service (back-end)
- Active/Active High Availability Support
- Enhanced support for PowerVC

OpenStack Features Roadmap

- Full REST v2 compliance (in review)
- SafeMode Support (in review)
- Volume Group Support (in progress)
- ActiveDR Support
- NVMe-FC Cinder driver for FlashArray (pending os-brick support)
- Manila driver for FlashArray File Services (2H 2025)
- Cinder NFS Drivers for FlashBlade and FlashArray

VMware to OpenStack Migration Options

- Utilise vVOLs
 - convert VMDK to vVOL
 - clone vVOL to full volume
 - import volume into OpenStack using `cinder manage`
- Use `qemu-img`
 - convert VMDK to QCOW2 using `qemu-img convert`
 - import to OpenStack as image using `openstack image create`
- Use V2V toolset (incl. backup/restore)
 - eg: Linux `virt-v2v`, Trilio, Commvault

DEMO TIME...

OpenStack is simple

```
[fsl-1]
volume_backend_name = ls-1
volume_driver = cinder.volume.drivers.pure.PureVCSIDriver
san_ip = 10.21.200.154
san_login = root
san_password = 533hh6-236f-261a-e0a2-6a0846cb122
use_chap_auth = False
pure_multipath_for_image_xfer = True
pure_msadlocate_on_delete = True
[oslo.middleware]
enable_proxy_headers_parsing = True

[nova]
region_name = RegionOne
remembered_servers = localhost:1211
osfile = /opt/stack/data/cs-bundle.pem
project_domain_name = Default
project_name = service
password_min_length = 10
password_max_length = 16
username = nova
auth_url = https://10.21.200.154/identity
interface = public
auth_type = password

[coordination]
backend_url = etcd3+http://10.21.200.154:2379/api_version=v3
[neutron_policy]


```

The Pure Storage Management UI interface. On the left, a sidebar menu includes: Storage, Storage, Protection, Analytics, Performance, Capacity, Health, and Settings. The main area has tabs for Array, Hosts, Volumes, Pools, File Systems, and Policies. The Volumes tab is currently selected. Below the tabs, there are two sections: 'Volumes' and 'Volume Groups'. Each section has a search bar, a table header (Name, Size, Virtual, Status, Snapshot, Dedupe), and a 'Create' button.

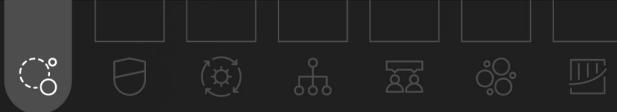
The OpenStack Horizon interface. The top navigation bar shows 'Admin' and 'Compute' selected. The main content area is titled 'Volumes' under the 'Admin' section. A table lists columns: Project, Host, Name, Size, Status, Group, Type, Attached To, Bootable, Encrypted, and Actions. A message at the bottom of the table says 'No items to display.'

Automation



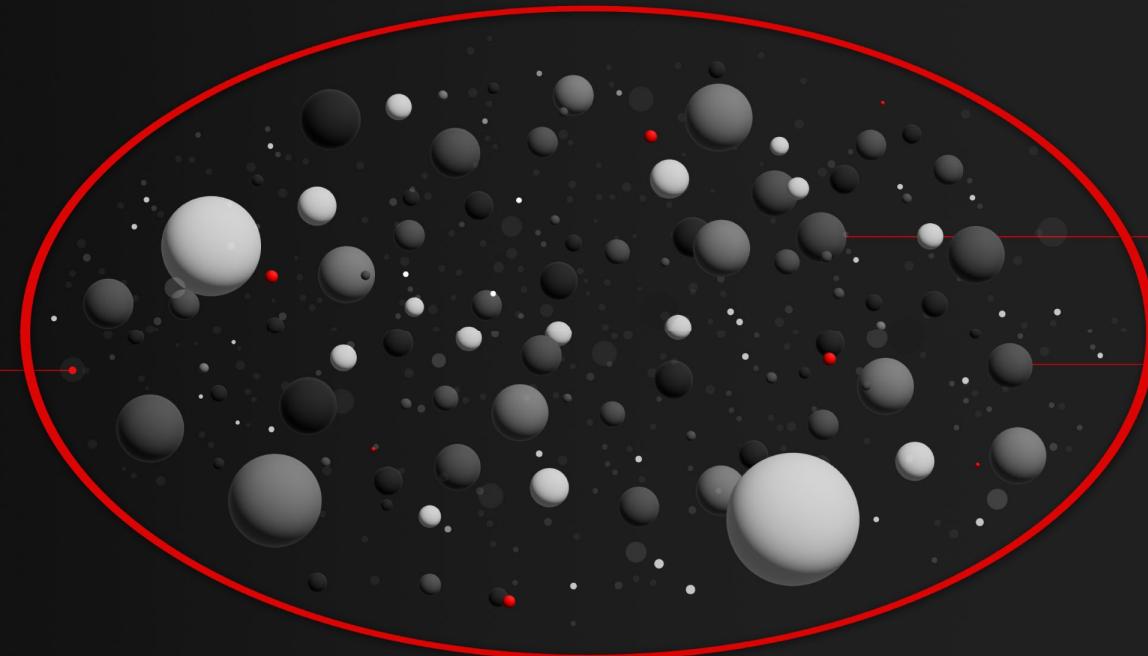


Ansible



#1

IT automation
community in
the world

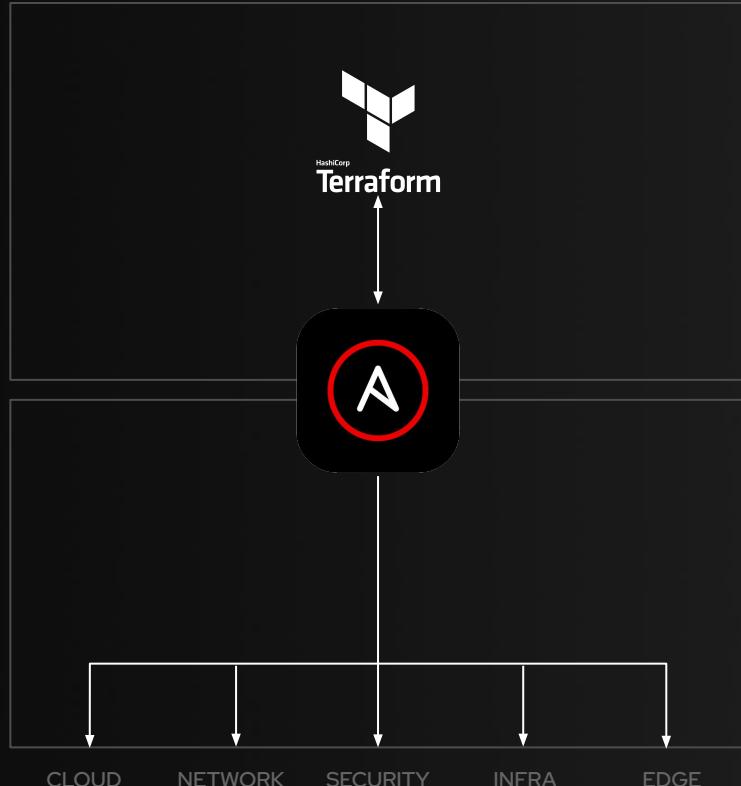


55,000+

GitHub stars

3,550+

Active open source
contributors



Everything as code
Everything is possible

► Infrastructure and configuration as code

- Ansible creates infrastructure manifests and triggers the provisioning and deprovisioning of infrastructure.
- Dynamic inventory management provides access to newly provisioned infrastructure without manual intervention.
- Combine tool chains to deliver infrastructure as code and configuration as code. Allowing for complete management of infrastructure life cycles with post-provisioning tasks.

One tool to manage multiple storage platforms

Day 0/1 deployment through Day 2 maintenance

Manage users, policies,
volumes, snapshots, etc. with
certified content collections

Reduce system inconsistencies and
cost by optimizing storage
management

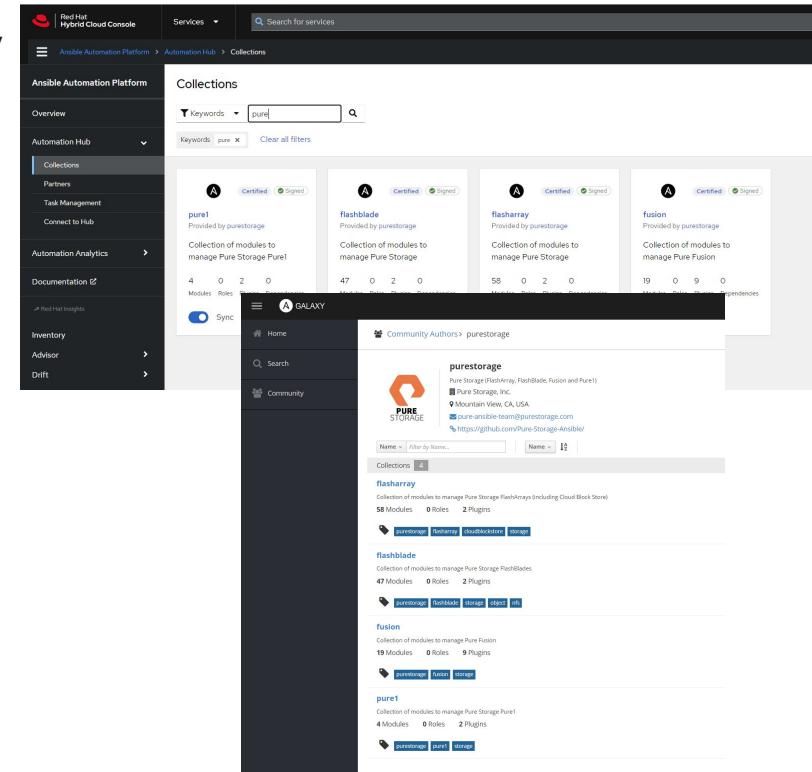
Configure
directory
services

Simplify provisioning



WHAT DOES PURE PROVIDE

- Modules and Collections Fully Certified by Red Hat
- Users can log support calls directly with Red Hat or Pure
- ... or raise [GitHub issues](#)
- Available through Collections
 - [Ansible Galaxy](#)
 - [Ansible Automation Hub](#) (RH account needed)
 - Containerized Execution Environments



The image shows two screenshots of the Ansible ecosystem. The top screenshot is the 'Automation Hub' collections page, where a search for 'pure' has returned four results: 'pure1', 'flashblade', 'flasharray', and 'fusion'. Each result shows a brief description, the number of modules, roles, and plugins, and a 'Certified' and 'Signed' badge. The bottom screenshot is the 'purestorage' collection page on Ansible Galaxy, which includes sections for 'purestorage' (Cloud Block Store), 'flasharray' (Cloud Block Storage), 'flashblade' (Pure Storage FlashBlade), and 'fusion' (Pure Fusion). Each section provides a brief description, the number of modules, roles, and plugins, and a list of tags.

Collections Fully Certified by Red Hat

- 117 modules and counting



Users can log support calls directly with Red Hat or Pure

- ... or raise [GitHub issues](#)

FA, FB and Fusion* included in latest Ansible versions

A screenshot of the Red Hat Hybrid Cloud Console interface. The left sidebar shows navigation options like Overview, Collections (which is selected), Partners, Repo Management, Connect to Hub, Automation Analytics, Documentation, Inventory, Advisor, and Drift. The main content area is titled "Collections" and shows search results for "pure". It displays four certified collections: "flasharray" (Provided by Pure Storage), "pure1" (Provided by Pure Storage), "flashblade" (Provided by Pure Storage), and "fusion" (Provided by Pure Storage). Each card shows the number of modules, roles, plugins, and dependencies. Sync buttons are present for each collection.



Ansible

PURE INFO MODULES

USED TO BE CALLED FACTS MODULES

Everything you could ever want to know about your platform.

Get all the facts or just a subset (sometimes much faster response)

purefa_info, purefb_info & pure1_info

FlashArray subsets available*:

all, minimum, config, performance, capacity, network, subnet, interfaces, hgroups, pgroups, hosts, admins, volumes, snapshots, replication, vgroups, offload, apps, arrays, certs, kmip, clients, policies, dir_snaps, filesystems, virtual_machines, hosts_balance and subscriptions

FlashBlade subsets available*:

all, minimum, config, performance, capacity, network, subnets, lags, filesystems, snapshots, buckets, replication, policies, arrays, accounts, admins, ad, kerberos and drives

Pure1 subsets available*:

all, minimum, appliances, subscriptions, contracts, environmental and invoices

Use the created {dict} to **set facts** to use later in the playbook and to do clever things with Jinja2

*Not a comprehensive list
Uncomplicate Data Storage, Forever



DEFINING YOUR TARGET PLATFORM

TELLING MODULES WHICH FA OR FB TO TALK TO...

- All modules need to have **fa_url** / **fb_url** and **api_token** provided as in the previous examples, except...
- Define shell environmental variables:
 - PUREFA_URL=10.21.200.210
 - PUREFA_API=c6033033-fe69-2515-a9e8-966bb7fe4b40
 - PUREFB_URL=10.21.200.5
 - PUREFB_API=T-9f276a18-50ab-446e-8a0c-666a3529a1b6
- These will be applied to **ALL** purefa or purefb modules, unless they have local definitions in the task.
- Or, use variables in the playbook, or as part of a variables file, or even pass in through the ansible command line.



PURE1 AS A TARGET

TELLING MODULES HOW TO TALK TO PURE1...

- All **pure1** modules need to have **app_id** and **key_file** provided with **password**, if set, except...
- Define shell environmental variables:
 - `PURE1_APP_ID=pure1:apikey:bRyL4DfW2KaItKMj`
 - `PURE1_PRIVATE_KEY_FILE="/directory/file.pem"`
 - `PURE1_PRIVATE_PASSWORD=password`
- These will be applied to **ALL** **pure1** modules, unless they have local definitions in the task.
- Or, use variables in the playbook, or as part of a variables file, or even pass in through the ansible command line.



ANSIBLE FACTS FOR EXTERNAL STORAGE

ADDED BY PURE STORAGE

- The following core facts for hosts are available to help automate storage provisioning when using `gather_facts: true` in your playbook or role
 - `ansible_iscsi_iqn` – provides the host iSCSI IQN
 - `ansible_hostnqn` – provide the host NVMe NQN
 - `ansible_fibre_channel_wwn` – provides all the host fibre channel WWNs, specifically the WWPNs
- Examples would look like this:

```
"ansible_iscsi_iqn": "iqn.1994-05.com.redhat:58acc9c1b3ca"  
"ansible_hostnqn": "nqn.2014-08.com.example:nvme:nvm-subsystem-sn-d78432"  
"ansible_fibre_channel_wwn": [  
    "21000024ff52a9ba",  
    "21000024ff52a9bb"  
]
```

- Coincidentally these are formatted as required by the Pure FlashArray module

SIMPLE EXAMPLES

USING COLLECTIONS

- Create a FlashArray volume and attach to the localhost:
- Create a FlashBlade NFS v4 filesystem with quotas, hard limit and snapshot support:

```
- name:  
  hosts: localhost  
  gather_facts: true  
  tasks:  
    - name: Create a volume  
      purestorage.flasharray.purefa_volume:  
        name: testvolume  
        size: 13G  
        bw_qos: 25M  
        iops_qos: 200K  
        fa_url: <FlashArray Mgmt VIP>  
        api_token: <API token>  
  
    - name: Create host and attach volume  
      purestorage.flasharray.purefa_host:  
        name: "{{ ansible_hostname }}"  
        iqn: "{{ ansible_iscsi_iqn }}"  
        volume: testvolume  
        fa_url: <FlashArray Mgmt VIP>  
        api_token: <API token>
```

```
- name:  
  hosts: localhost  
  tasks:  
    - name: Create a filesystem  
      purestorage.flashblade.purefb_fs:  
        name: testshare  
        size: 1T  
        nfsv4: True  
        user_quota: 10k  
        group_quota: 100k  
        hard_limit: True  
        snapshot: True  
        nfs_rules: "10.10.28.78/32(rw,no_root_squash)"  
        fb_url: <FlashBlade Mgmt VIP>  
        api_token: <API token>
```



FURTHER EXAMPLES

USING COLLECTIONS

- Connect two FlashArrays in ActiveCluster and create a stretched pod:

```
- name:  
  hosts: localhost  
  tasks:  
    - name: Connect FlashArrays  
      purestorage.flasharray.purefa_connect:  
        connection: sync  
        target_url: <Target FA Mgmt VIP>  
        target_api: <Target FA API token>  
        fa_url: <Source FA Mgmt VIP>  
        api_token: <Source FA API token>  
  
    - name: Create Active Cluster pod  
      purestorage.flasharray.purefa_pod:  
        name: pod_1  
        fa_url: <FlashArray Mgmt VIP>  
        api_token: <FlashArray API token>  
  
    - name: Stretch Active Cluster pod  
      purestorage.flasharray.purefa_pod:  
        name: pod_1  
        stretch: <target array name>  
        fa_url: <FlashArray Mgmt VIP>  
        api_token: <FlashArray API token>
```

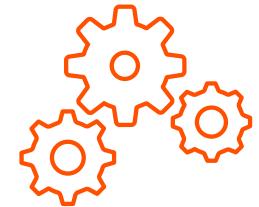
- Connect two FlashBlades and create a filesystem replica:

```
- name:  
  hosts: localhost  
  tasks:  
    - name: Connect FlashBlades  
      purestorage.flashblade.purefb_connect:  
        target_url: <Target FB Mgmt VIP>  
        target_api: <Target FB API token>  
        fb_url: <Source FB Mgmt VIP>  
        api_token: <Source FA API token>  
  
    - name: Create File System Replica  
      purestorage.flashblade.purefb_fs_replica:  
        name: fs_1  
        fb_url: <FlashBlade Mgmt VIP>  
        api_token: <FlashBlade API token>
```



USE CASES WITH PURE + ANSIBLE...

DAILY PROVISIONING, CONFIGURATION MANAGEMENT, ETC...



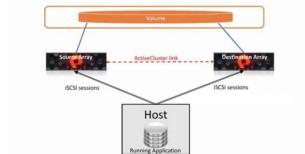
- Configure new arrays to the corporate standards
- Manage infrastructure configuration drift on fleets of arrays
- Update global configuration settings across multiple arrays
- Perform repetitive tasks, eg clone DB volumes for development work
- Perform DR tasks in a consistent and reproducible fashion
- Deploy hosts with a consistent external storage config

**IF YOU CAN DO IT ON THE GUI OR CLI,
YOU CAN DO IT WITH ANSIBLE**

REAL WORLD EXAMPLES...



[Manage Infrastructure Drift](#)



[Live Application Data Migration](#)



[Application Disaster Recovery in Kubernetes*](#)



[Load Balancing FlashArrays](#)



[Oracle Multi-Database Refresh](#)

Red Hat Solution Brief

Automate your storage environment with Red Hat and Pure Storage



Data management is a growing concern

87% of IT leaders believe that managing unstructured data growth is a top priority.

68% of organizations spend more than **30%** of their IT budget on data storage, backups, and disaster recovery.

Data growth brings new IT challenges

Modern enterprise IT organizations manage enormous amounts of data. In fact, more than 50% of enterprises have more than 5 petabytes (PB) of data.¹ This quantity of data can make it difficult for IT teams to operate efficiently and respond rapidly to issues, even in highly available and security-focused data storage environments.

As a result, many organizations want to streamline data resilience and security processes and simplify high-volume, repetitive tasks like provisioning and backup. Red Hat and Pure Storage offer solutions to help you increase the efficiency and accuracy of day-to-day operations and improve response time.

Build automated data storage workflows with Red Hat and Pure Storage

Red Hat and Pure Storage combine advanced automation and data storage platforms into integrated, streamlined solutions. As a foundation for building and scaling automation across an organization, Red Hat® Ansible® Automation Platform includes all of the tools you need to implement enterprise-wide automation in hybrid cloud environments.

Pure Storage delivers a modern data experience with flash-based storage and policy-driven services that let you run an automated, Storage-as-a-Service model across your storage platforms. With Pure Storage, you can build an all-flash datacenter, based on FlashBlade, FlashArray, and Pure Fusion's storage-as-code capabilities to reduce the complexity and expense of managing your environment.

Ansible Pure Storage Certified Content collections—predefined automation content developed, tested, and maintained by Red Hat and Pure Storage—help you automate more rapidly. Published to Ansible automation hub, these collections integrate Ansible Automation Platform and Pure Storage products, including FlashArray, FlashBlade, Pure1, and Fusion. You can perform nearly every operation in the Pure Storage command line and graphical user interfaces (GUIs) directly from Ansible Automation Platform as part of advanced automation workflows for managing your entire data storage environment.

Snapshot production databases

When integrating and testing new application features, developers need to work with the latest production data. IT staff typically use manual processes to copy production databases for developers to use. However, these databases can quickly diverge, causing issues when deploying new features into production. Ansible Automation Platform and Pure Storage let you build automated workflows that clone production databases each morning and delete old copies late at night. You can keep your production database safe, provide up-to-date copies for developers, and manage storage—all without manual intervention.

Ansible-foo

Troubleshooting Ansible
for Beginners

- 01** **Latest Versions**
Ensure you are at the latest versions of Ansible*, Collections and Pure SDKs
- 02** **Multiple Versions**
Make sure you don't have multiple Collections in different locations - `ansible-galaxy collection list`
- 03** **Check your YAML**
YAML syntax is finicky - test with a YAML lint validator online
- 04** **Verbosity is your friend**
Rerun failures with `-vvvv` to get extended traceback information
- 05** **Devil is in the Detail**
When raising support calls or GitHub issues, provide versions, logs and example playbooks



VERSIONS ARE IMPORTANT

GETTING THE DETAILS

- Ansible version

```
$ ansible --version
ansible [core 2.15.3]
  config file = None
  configured module search path = ['/home/simon/.ansible/plugins/modules',
 '/usr/share/ansible/plugins/modules']
    ansible python module location = /home/simon/.local/lib/python3.10/site-packages/ansible
    ansible collection location = /home/simon/.ansible/collections:/usr/share/ansible/collections
    executable location = /home/simon/.local/bin/ansible
    python version = 3.10.12 (main, Jun 11 2023, 05:26:28) [GCC 11.4.0] (/usr/bin/python3)
    jinja version = 3.1.2
    libyaml = True
```

- Ansible Collections versions

```
$ ansible-galaxy collection list | grep purestorage
purestorage.flasharray      1.21.0
purestorage.flashblade      1.13.1
purestorage.pure1           1.3.0
purestorage.flashblade      1.12.1
```



Notice two different versions

- SDK versions

```
$ pip list | grep pur
purestorage                  1.19.0
purity-fb                    1.12.3
py-pure-client                1.41.1
```



YAML IS “FUN”

SPOT THE DIFFERENCE

```
simon@simon-X1: $ ansible-playbook pure1.yaml
[WARNING]: No inventory was parsed, only implicit localhost is available
[WARNING]: provided hosts list is empty, only localhost is available. Note that the implicit localhost does not match 'all'
ERROR! We were unable to read either as JSON nor YAML, these are the errors we got from each:
JSON: Expecting value: line 1 column 1 (char 0)

Syntax Error while loading YAML.
  mapping values are not allowed in this context

The error appears to be in '/home/simon/pure1.yaml': line 8, column 17, but may
be elsewhere in the file depending on the exact syntax problem.

The offending line appears to be:

- name: Collect volumes information
  pure1_audits:
    ^ here
```

```
1 ---  
2   - name: Pure Storage storage module examples  
3     hosts: localhost  
4     gather_facts: no  
5     collections:  
6       - purestorage.pure1  
7         tasks:  
8           - name: Collect volumes information  
9             pure1_audits:  
10               app_id: pure1:spikekey:bRyL4DfW2KaItkMj  
11               key_file: /home/simon/puretec.pem  
12             - name: Collect information of the source Pure Storage FlashArray  
13               pure1_info:  
14                 gather_subset: subscriptions  
15                 app_id: pure1:spikekey:bRyL4DfW2KaItkMj  
16                 key_file: /home/simon/puretec.pem  
17             - pause: none  
18             - name: Collect information of the source Pure Storage FlashArray  
19               pure1_alerts:  
20                 severity: info
```

Below the code editor are two buttons: "Go" and "Reformat (strips comments) Resolve aliases". A red box highlights the warning message at the bottom:

Nested mappings are not allowed in compact mappings at line 8, column 12
Implicit keys need to be on a single line at line 8, column 12
All sequence items must start at the same column at line 12, column 5

```
1 ---  
2   - name: Pure Storage storage module examples  
3     hosts: localhost  
4     gather_facts: no  
5     collections:  
6       - purestorage.pure1  
7         tasks:  
8           - name: Collect volumes information  
9             pure1_audits:  
10               app_id: pure1:spikekey:bRyL4DfW2KaItkMj  
11               key_file: /home/simon/puretec.pem  
12             - name: Collect information of the source Pure Storage FlashArray  
13               pure1_info:  
14                 gather_subset: subscriptions  
15                 app_id: pure1:spikekey:bRyL4DfW2KaItkMj  
16                 key_file: /home/simon/puretec.pem  
17             - pause: none  
18             - name: Collect information of the source Pure Storage FlashArray  
19               pure1_alerts:  
20                 severity: info
```

Below the code editor are two buttons: "Go" and "Reformat (strips comments) Resolve aliases". A green box highlights the success message at the bottom:

Valid YAML!

USING - VVVV

ansible-playbook -vvvv <playbook>

```
TASK [aaa]
*****
task path: /etc/ansible/roles/purefa_snap_repl_kddia.yml:4
<127.0.0.1> ESTABLISH LOCAL CONNECTION FOR USER: root
<127.0.0.1> EXEC /bin/sh -c 'echo ~root && sleep 0'
<127.0.0.1> EXEC /bin/sh -c '( umask 77 && mkdir -p ` echo /root/.ansible/tmp `&& mkdir ` echo /root/.ansible/tmp/ansible-tmp-1695703335.1349988-35950-135452677291121 ` && echo /root/.ansible/tmp/ansible-tmp-1695703335.1349988-35950-135452677291121 ) && sleep 0'
Using module file /usr/local/lib/python3.9/site-packages/ansible_collections/purestorage/flasharray/plugins/modules/purefa_snap.py
<127.0.0.1> PUT /root/.ansible/tmp/ansible-local-35893ae91kwvl/tmpmgw8xmbg TO
/root/.ansible/tmp/ansible-tmp-1695703335.1349988-35950-135452677291121/Ansiballz_purefa_snap.py
<127.0.0.1> EXEC /bin/sh -c 'chmod u+x /root/.ansible/tmp/ansible-tmp-1695703335.1349988-35950-135452677291121/
/root/.ansible/tmp/ansible-tmp-1695703335.1349988-35950-135452677291121/Ansiballz_purefa_snap.py && sleep 0'
<127.0.0.1> EXEC /bin/sh -c '/usr/bin/python3 /root/.ansible/tmp/ansible-tmp-1695703335.1349988-35950-135452677291121/Ansiballz_purefa_snap.py && sleep 0'
<127.0.0.1> EXEC /bin/sh -c 'rm -f -r /root/.ansible/tmp/ansible-tmp-1695703335.1349988-35950-135452677291121/ > /dev/null 2>&1 && sleep 0'
fatal: [localhost]: FAILED! => {
    "changed": false,
    "module_stderr": "/usr/lib/python3.9/site-packages/urllib3/connectionpool.py:1013: InsecureRequestWarning: Unverified HTTPS request is being made to host '10.229.192.60'. Adding certificate verification is strongly advised. See: \nhttps://urllib3.readthedocs.io/en/latest/advanced-usage.html#ssl-warnings\n",
    "module_stdout": "warnings.warn('/n/usr/lib/python3.9/site-packages/urllib3/connectionpool.py:1013: InsecureRequestWarning: Unverified HTTPS request is being made to host '10.229.192.60'. Adding certificate verification is strongly advised. See: \nhttps://urllib3.readthedocs.io/en/latest/advanced-usage.html#ssl-warnings\n",
    "msg": "warnings.warn('/n/usr/lib/python3.9/site-packages/urllib3/connectionpool.py:1013: InsecureRequestWarning: Unverified HTTPS request is being made to host '10.229.192.60'. Adding certificate verification is strongly advised. See: \nhttps://urllib3.readthedocs.io/en/latest/advanced-usage.html#ssl-warnings\n",
    "rc": 1
}
```



DEMO TIME...

Some simple Ansible workflows

```
[root@snl-pool-c07-15 ~]#
```

```
Every 1.0s: multipath -ll      Mon Feb 12 16:28:43 2024
Every 1.0s: iscscliadm -m session  Mon Feb 12 16:28:44 2024
iscscliadm: No active sessions.
```

The screenshot shows the Pure Storage Management UI interface. On the left, there is a navigation sidebar with icons for Dashboard, Storage (selected), Protection, Analytics, Performance, Capacity, and Replication, followed by Health and Settings. Below the sidebar are links for Help, End User Agreement, Terms, and Log Out, along with a user icon and name (Alice TA1). The main content area has two tabs: 'Hosts' (selected) and 'Volumes'. The 'Hosts' tab displays a table with one row: 'Hosts' (Value: 600, Data Reduction: 0TB, Unique: 0TB, Snapshot: 0TB, Total: 0TB). It also includes sections for 'Host Groups' and a search bar. The 'Volumes' tab displays a table with one row: 'Volume' (Value: 1, Data Reduction: 0TB, Unique: 0TB, Snapshot: 0TB, Total: 0TB). It includes sections for 'Volume Groups', 'Specs', 'QoS', and 'Details', along with a search bar. Both tabs have tabs for Array, Hosts, Volumes, Pools, File Systems, and Policies.

Observability



Pure and O11y

Meet you where your data
is...

- Whatever platform you choose
- Integrations based on standards
 - OpenMetrics
 - OpenTelemetry



- Open Metrics Exporters for FlashArray and FlashBlade
 - Runs in Docker container
 - Best Efforts support only
- Integrates into all major O11y platforms
- Get them for free...
 - <https://github.com/PureStorage-OpenConnect/pure-fa-openmetrics-exporter>
 - <https://github.com/PureStorage-OpenConnect/pure-fb-openmetrics-exporter>
- Future: Integration directly into Purity//FA 6.6.8 and //FB
 - Full support at this point

OpenMetrics Exporter



- Open Telemetry Receivers for FlashArray and FlashBlade
- Merged into otel-collector-contrib repo
- Can be used by any OTel compliant Observability toolset
 - Already merged into SumoLogic and OTel Collector Contrib distributions
- More details...
 - <https://github.com/open-telemetry/opentelemetry-collector-contrib/tree/main/receiver/purefareceiver>
 - <https://github.com/open-telemetry/opentelemetry-collector-contrib/tree/main/receiver/purefbreceiver>

OpenTelemetry Receivers



Current OME Integrations

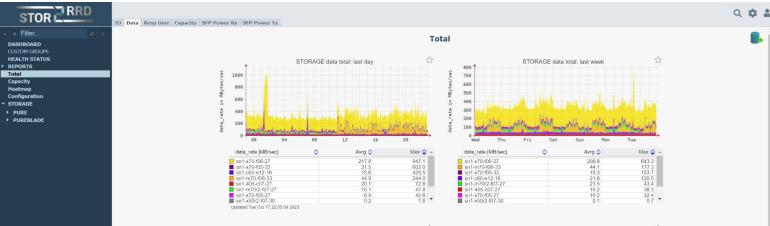
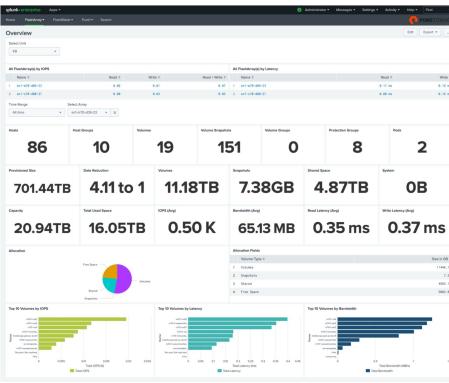
All use the same source data feeds from OpenMetrics Exporters

	Cloud-based (SaaS)	On-Prem (Darksite)
	✗	✓
	✓	✓
	✓	✗
	✓	✗

Using proprietary
connection agents

Other Monitoring Tools

splunk>enterprise



DEMO TIME...

Let's see these current integrations in action

